

Kanu Madhok

952-303-1045 • madhok.kanu@gmail.com • github.com/kmadhok • linkedin.com/in/kanu-madhok

PROFESSIONAL EXPERIENCE

Walmart Data Ventures

Chicago, IL · Sep 2024 – Present

Senior Data Analyst

- Designed and shipped a self-service analytics agent over 67 BigQuery tables with sub-agent orchestration and context-isolated validation (Context Researcher → SQL Drafter → Validator → Devil's Advocate); 8 deterministic SQL rules, 23 golden-query test cases. Approved by Sr. Director and product leadership and currently the only AI skill in active use by business teams on Data Ventures. Eliminated the analyst as the bottleneck for routine asks.
- Built an autonomous Jira ticket-resolution agent: 6-gate triage, MCP-driven context gathering, FAISS+BM25 retrieval over 11,000 historical SQL queries, self-healing execution with a 10-retry budget. Closed 400+ data requests across 8 categories over 15 months — cut turnaround from 30–60 minutes per request (and up to a full day for complex ones) to under 10 minutes each. Stakeholders never knew an agent was answering them.
- Built a hybrid orchestrator on top of the analytics agent. Routes business questions through two backends in parallel: a CubeJS semantic-layer skill (co-built with the DS and product team) and the context-engineered agent itself. Compares outputs, reasons about disagreements, and feeds discrepancies back to refine the semantic layer's definitions over time.
- Built a headless KPI monitor on a 6-hour cron that detects threshold breaches, runs root-cause analysis, and posts Slack alerts; state persists to BigQuery so the agent has no LLM-memory reliance. Designed a v2 architecture as a 3-tier event-driven platform (reactive events, threshold KPIs, PM-owned themes) self-configuring via a BigQuery event registry. New event types ship without code deploys.
- Replaced ~20 distinct recruitment workflows I'd been running by hand 1–2× weekly with end-to-end automation on a 68-table BigQuery platform. Stratified-sampling pipelines recruited 280,000+ panelists across 29 categories with zero duplicates.

University of Chicago — Data Science Institute

Chicago, IL · Jan – Sep 2024

Graduate Researcher, Data & Democracy Initiative

- Persona-driven multi-LLM (OpenAI + Gemini) chatbot for a 500+ participant donation experiment; controlled A/B tests persuaded 15% to donate to an opposing cause.

FTI Consulting — UChicago MSADS Capstone

Chicago, IL · Jan – Aug 2024

- RAG-to-proposal generator deployed on AWS EC2; cut proposal time ~30%; Best in Show — UChicago MSADS Capstone.

Innovare (EdTech Startup)

Chicago, IL · Oct 2022 – Mar 2023

Data Scientist

- Logistic-regression risk model + Looker dashboards for 10+ districts (5,000+ students); retention +7%, failures –15%.

SELECTED PROJECT

- Live deployed text-to-SQL copilot (Streamlit + Gemini 2.5 + RAG over historical queries, BigQuery execution, auto-visualization): live demo — <https://sql-rag-frontend-simple-481433773942.us-central1.run.app/>

SKILLS

Agent stack: MCP (Model Context Protocol) server design, sub-agent orchestration, context engineering, prompt engineering, RAG (FAISS + BM25 + RRF), ChromaDB semantic memory, multi-model routing across providers, headless cron agents, event-driven architectures

LLMs: Claude (Opus / Sonnet / Haiku via Wibey), Gemini (2.5 Pro / Flash via Vertex AI), OpenAI

APIs / triggers / glue: FastAPI, Cloud Run, BigQuery scheduled queries, Slack webhooks, cron, Playwright SSO, JSONL state stores

Data: BigQuery, dbt, SQL, Python, PySpark, Airflow, Databricks, MySQL

Viz: Power BI (DAX, XMLA TMSL), Tableau, LookML, Streamlit, Matplotlib, React Recharts

EDUCATION

The University of Chicago — M.S., Applied Data Science (GPA 4.0)

Chicago, IL

Loyola University Chicago — B.S., Information Systems & Economics (GPA 3.5)

Chicago, IL