

Kanu Madhok

952-303-1045 • madhok.kanu@gmail.com • github.com/kmadhok • linkedin.com/in/kanu-madhok

PROFESSIONAL EXPERIENCE

Walmart Data Ventures

Chicago, IL · Sep 2024 – Present

Senior Data Analyst

- Designed and shipped to production a self-service analytics agent over 67 BigQuery tables with sub-agent orchestration and human-in-the-loop validation; 8 deterministic SQL rules and a 23-case golden-query evaluation suite. Partnered with Product, Data Science, and business stakeholders to translate ambiguous business questions into agent-ready specs. Approved by Sr. Director and product leadership and currently the only AI skill in active use by business teams on Data Ventures.
- Built and deployed to production a headless KPI monitor on a 6-hour cron that detects threshold breaches, runs root-cause analysis, and posts Slack alerts; end-to-end observability via decision logging and BigQuery state persistence. Designed a v2 architecture as a 3-tier event-driven platform (reactive events, threshold KPIs, PM-owned themes) self-configuring via a BigQuery event registry — new event types ship without code deploys.
- Built and shipped to production an autonomous Jira ticket-resolution agent: 6-gate triage, MCP-driven context gathering, agentic RAG with FAISS + BM25 retrieval over 11,000 historical SQL queries, and self-healing execution with a 10-retry budget. Closed 400+ data requests across 8 categories over 15 months — cut routine turnaround from 30–60 minutes to under 10 minutes each.
- Evolved the analytics skill into a registry-governed semantic layer over BigQuery: a single LLM call maps the question to strict JSON intent (the model never writes SQL); a deterministic gate validates and canonicalizes intent before a deterministic compiler emits the SQL — same question, same answer, every time. Definition changes pass a human review gate before entering the registry.
- Replaced ~20 distinct recruitment workflows with end-to-end automation on a 68-table BigQuery platform; stratified-sampling pipelines recruited 280,000+ panelists across 29 categories with zero duplicates. A recurring 3-hour task now takes ~3 minutes of setup.

University of Chicago — Data Science Institute

Chicago, IL · Jan – Sep 2024

Graduate Researcher, Data & Democracy Initiative

- Persona-driven multi-LLM (OpenAI + Gemini) chatbot for a 500+ participant donation experiment; controlled A/B tests persuaded 15% to donate to an opposing cause.

FTI Consulting — UChicago MSADS Capstone

Chicago, IL · Jan – Aug 2024

- RAG-to-proposal generator deployed on AWS EC2; cut proposal time ~30%; Best in Show — UChicago MSADS Capstone.

Innovare (EdTech Startup)

Chicago, IL · Oct 2022 – Mar 2023

Data Scientist

- Logistic-regression risk model + Looker dashboards for 10+ districts (5,000+ students); retention +7%, failures –15%.

SELECTED PROJECT

- Live deployed text-to-SQL copilot (Streamlit + Gemini 2.5 + LangChain-orchestrated RAG over historical queries, BigQuery execution, auto-visualization): live demo — <https://sql-rag-frontend-simple-481433773942.us-central1.run.app/>

SKILLS

Agent stack: sub-agent orchestration, context engineering, prompt engineering, RAG (embeddings + FAISS + BM25 + RRF), LangChain, ChromaDB semantic memory, evaluation suites (golden sets, human-in-the-loop validation), MCP server design, headless cron agents, event-driven architectures, decision logging / observability

Data: Python, SQL, BigQuery, dbt, PySpark, Airflow, Databricks, MySQL

LLMs: Claude (Opus / Sonnet / Haiku via Wibey), Gemini (2.5 Pro / Flash via Vertex AI), OpenAI

AI-assisted development: Claude Code, Cursor

APIs / triggers / glue: FastAPI, Cloud Run, BigQuery scheduled queries, Slack webhooks, cron, Playwright SSO, JSONL state stores, CI/CD (pytest, version control)

Viz / reporting: Streamlit, Power BI (DAX, XMLA TMSL), Tableau, LookML, Matplotlib, React Recharts

EDUCATION

The University of Chicago — M.S., Applied Data Science (GPA 4.0)

Chicago, IL

Loyola University Chicago — B.S., Information Systems & Economics (GPA 3.5)

Chicago, IL