

# Kanu Madhok

952-303-1045 • madhok.kanu@gmail.com • [github.com/kmadhok](https://github.com/kmadhok) • [linkedin.com/in/kanu-madhok](https://linkedin.com/in/kanu-madhok)

## PROFESSIONAL EXPERIENCE

### Walmart Data Ventures

Chicago, IL · Sep 2024 – Present

Senior Data Analyst

- Designed and shipped to production a self-service analytics agent over 67 BigQuery tables with sub-agent orchestration and human-in-the-loop validation (Context Researcher → SQL Drafter → Validator → Devil's Advocate); 8 deterministic SQL rules and a 23-case golden-query evaluation suite. Partnered with Product, Data Science, and business stakeholders to translate ambiguous business questions into agent-ready specs. Approved by Sr. Director and product leadership and currently the only AI skill in active use by business teams on Data Ventures. Eliminated the analyst as the bottleneck for routine asks.
- Built and deployed to production a headless KPI monitor on a 6-hour cron that detects threshold breaches, runs root-cause analysis, and posts Slack alerts; end-to-end observability via decision logging and BigQuery state persistence, so the agent has no LLM-memory reliance. Designed a v2 architecture as a 3-tier event-driven platform (reactive events, threshold KPIs, PM-owned themes) self-configuring via a BigQuery event registry — new event types ship without code deploys.
- Built and shipped to production an autonomous Jira ticket-resolution agent: 6-gate triage, MCP-driven context gathering, agentic RAG with FAISS + BM25 retrieval over 11,000 historical SQL queries, and self-healing execution with a 10-retry budget. Closed 400+ data requests across 8 categories over 15 months — cut routine turnaround from 30–60 minutes to under 10 minutes each (complex requests from up to a full day). Stakeholders never knew an agent was answering them.
- Replaced ~20 distinct recruitment workflows I had been running by hand 1–2× weekly with end-to-end automation on a 68-table BigQuery platform. Stratified-sampling pipelines recruited 280,000+ panelists across 29 categories with zero duplicates. A recurring 3-hour task now takes ~3 minutes of setup.
- Evolved the analytics skill into a registry-governed semantic layer over BigQuery: YAML definitions plus SQLite vocabulary snapshots define legal dimensions, values, and table wiring; a single LLM call maps the question to strict JSON intent (the model never writes SQL); a deterministic gate validates and canonicalizes intent before a deterministic compiler emits the SQL — same question, same answer, every time, with the only probabilistic step regression-testable. New or changed definitions pass a human review gate before entering the registry.
- Delivered 5 Power BI dashboards ranked top 1.4% of Walmart's 207,000 dashboards (1,064 views, 22 users); cut data-request turnaround from 1–24 hours to 5–10 minutes; Making a Difference Award.

### University of Chicago — Data Science Institute

Chicago, IL · Jan – Sep 2024

Graduate Researcher, Data & Democracy Initiative

- Persona-driven multi-LLM (OpenAI + Gemini) chatbot for a 500+ participant donation experiment; controlled A/B tests persuaded 15% to donate to an opposing cause.

### FTI Consulting — UChicago MSADS Capstone

Chicago, IL · Jan – Aug 2024

- RAG-to-proposal generator deployed on AWS EC2; cut proposal time ~30%; Best in Show — UChicago MSADS Capstone.

### Innovare (EdTech Startup)

Chicago, IL · Oct 2022 – Mar 2023

Data Scientist

- Logistic-regression risk model + Looker dashboards for 10+ districts (5,000+ students); retention +7%, failures –15%.

## SELECTED PROJECT

- Live deployed text-to-SQL copilot (Streamlit + Gemini 2.5 + RAG over historical queries, BigQuery execution, auto-visualization): live demo — <https://sql-rag-frontend-simple-481433773942.us-central1.run.app/>

## SKILLS

**Data:** BigQuery, dbt, SQL, Python, PySpark, Airflow, Databricks, MySQL

**Agent & evaluation stack:** MCP (Model Context Protocol) server design, sub-agent orchestration, context engineering, RAG (FAISS + BM25 + RRF), evaluation suites (golden sets, human-in-the-loop validation), headless cron agents, event-driven architectures, decision logging / observability

**LLMs:** Claude (Opus / Sonnet / Haiku via Wibey), Gemini (2.5 Pro / Flash via Vertex AI), OpenAI

**APIs / triggers / glue:** FastAPI, Cloud Run, BigQuery scheduled queries, Slack webhooks, cron, Playwright SSO, JSONL state stores, CI/CD (pytest, version control)

**Viz / reporting:** Power BI (DAX, XMLA TMSL), Tableau, LookML, Streamlit, Matplotlib, React Recharts

## EDUCATION

The University of Chicago — M.S., Applied Data Science (GPA 4.0)

Chicago, IL

Loyola University Chicago — B.S., Information Systems & Economics (GPA 3.5)

Chicago, IL